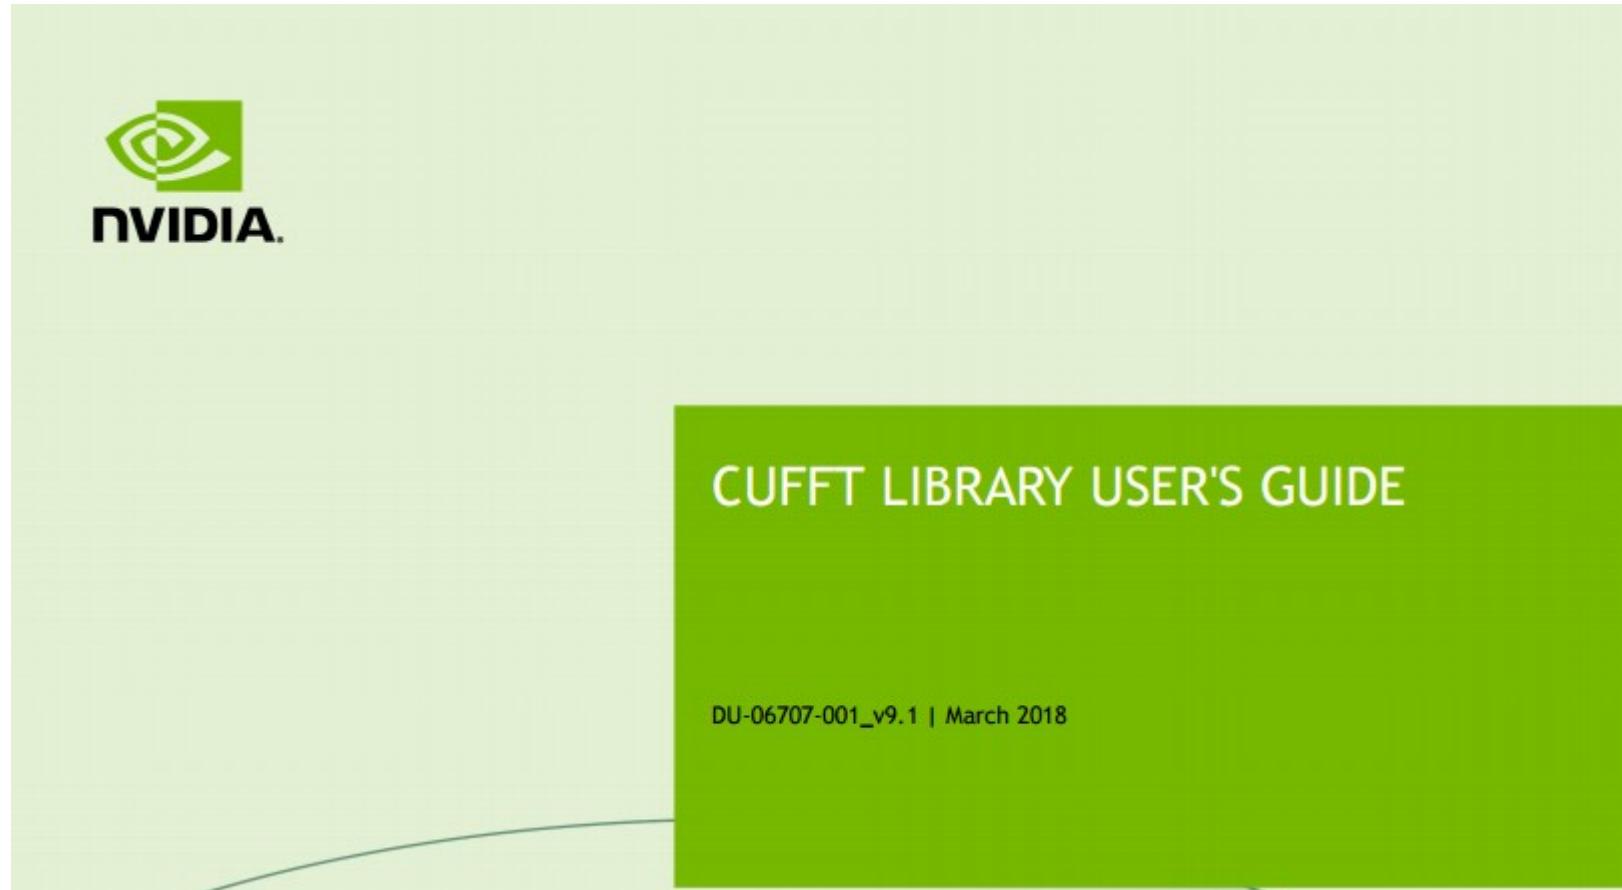


ICNPG 2023

Transformadas de Fourier



¿ Que es la CUFFT ?



The **FFT** is a divide-and-conquer algorithm for efficiently computing discrete Fourier transforms of complex or real-valued data sets. It is one of the most important and widely used numerical algorithms in computational physics and general signal processing. The CUFFT library provides a simple interface for computing FFTs on an NVIDIA GPU, which allows users to quickly leverage the floating-point power and parallelism of the GPU in a highly optimized and tested FFT library.

¿ Que es la CUFFT ?



This document describes cuFFT, the NVIDIA® CUDA™ Fast Fourier Transform (FFT) product. It consists of two separate libraries: cuFFT and cuFFTW. The cuFFT library is designed to provide high performance on NVIDIA GPUs. The cuFFTW library is provided as a porting tool to enable users of FFTW to start using NVIDIA GPUs with a minimum amount of effort.

The cuFFT product supports a wide range of FFT inputs and options efficiently on NVIDIA GPUs. This version of the cuFFT library supports the following features:

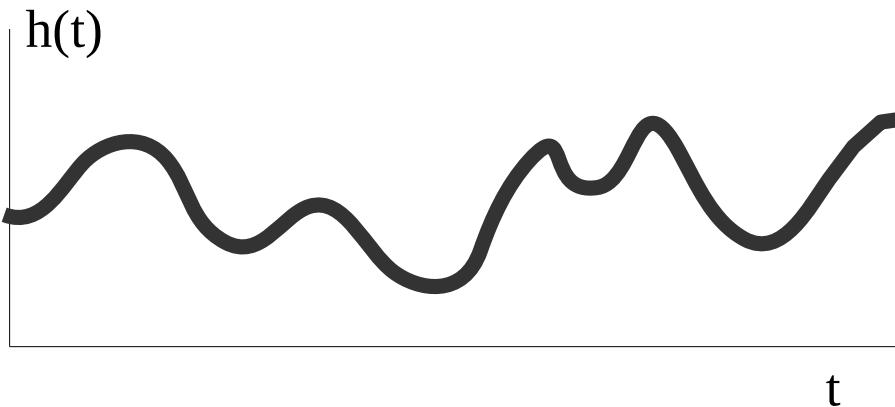
- ▶ Algorithms highly optimized for input sizes that can be written in the form $2^a \times 3^b \times 5^c \times 7^d$. In general the smaller the prime factor, the better the performance, i.e., powers of two are fastest.
- ▶ An $O(n\log n)$ algorithm for every input data size
- ▶ Half-precision (16-bit floating point), single-precision (32-bit floating point) and double-precision (64-bit floating point). Transforms of lower precision have higher performance.
- ▶ Complex and real-valued input and output. Real valued input or output require less computations and data than complex values and often have faster time to solution.
Types supported are:
 - ▶ C2C - Complex input to complex output
 - ▶ R2C - Real input to complex output
 - ▶ C2R - Symmetric complex input to real output
- ▶ 1D, 2D and 3D transforms
- ▶ Execution of multiple 1D, 2D and 3D transforms simultaneously. These batched transforms have higher performance than single transforms.
- ▶ In-place and out-of-place transforms modifica o no el array de entrada
- ▶ Arbitrary intra- and inter-dimension element strides (strided layout)

Por ej esto es necesario hacerlo para sacar el espectro de un audio a tiempo real

Si tenemos una señal muy larga y queremos hacer transformadas de a pedazos, podemos decirle como input. Esto es más eficiente que mandar en serie transformadas de señales pequeñas

Se pueden acceder a los datos de forma no continua, "pegando saltos" sin tener que reordenarlos

Repaso de Transformada de Fourier



Simetrías

If ...

$h(t)$ is real

$h(t)$ is imaginary

$h(t)$ is even

$h(t)$ is odd

$h(t)$ is real and even

$h(t)$ is real and odd

$h(t)$ is imaginary and even

$h(t)$ is imaginary and odd

then ...

$H(-f) = [H(f)]^*$

$H(-f) = -[H(f)]^*$

$H(-f) = H(f)$ [i.e., $H(f)$ is even]

$H(-f) = -H(f)$ [i.e., $H(f)$ is odd]

$H(f)$ is real and even

$H(f)$ is imaginary and odd

$H(f)$ is imaginary and even

$H(f)$ is real and odd

“Descomposición en senos y cosenos”

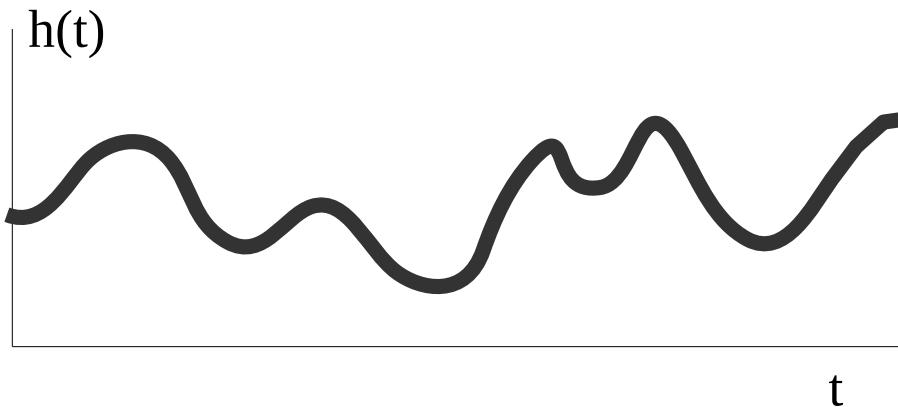
$$H(\omega) = \int_{-\infty}^{\infty} h(t)e^{i\omega t} dt$$

$$h(t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} H(\omega)e^{-i\omega t} d\omega$$

Ver Por ej:

Press et al, Numerical Recipes

Repaso de Transformada de Fourier



$$H(\omega) = \int_{-\infty}^{\infty} h(t)e^{i\omega t} dt$$
$$h(t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} H(\omega)e^{-i\omega t} d\omega$$

Esto es útil si tenemos que
reescalar nuestro problema

$$h(at) \iff \frac{1}{|a|} H\left(\frac{f}{a}\right)$$

time scaling

$$\frac{1}{|b|} h\left(\frac{t}{b}\right) \iff H(bf)$$

frequency scaling

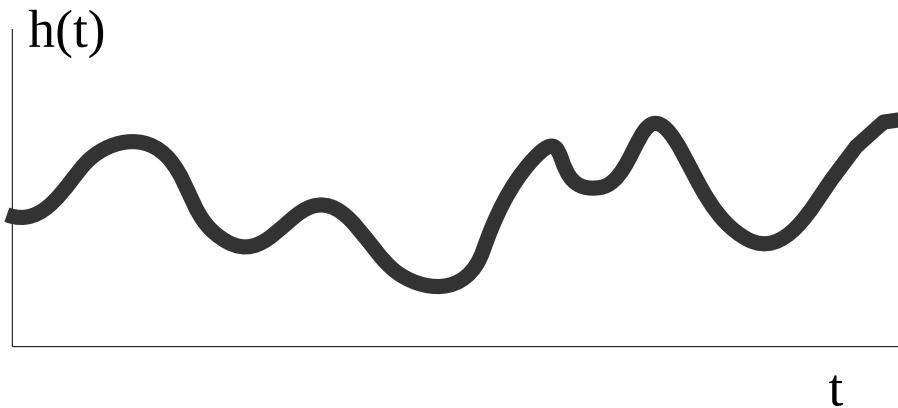
$$h(t - t_0) \iff H(f) e^{2\pi i f t_0}$$

time shifting

$$h(t) e^{-2\pi i f_0 t} \iff H(f - f_0)$$

frequency shifting

Repaso de Transformada de Fourier



“Descomposición en senos y cosenos”

$$H(\omega) = \int_{-\infty}^{\infty} h(t)e^{i\omega t} dt$$
$$h(t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} H(\omega)e^{-i\omega t} d\omega$$

Porque es tan importante la transformada de Fourier en física computacional ?

Convolution theorem

$$g * h \equiv \int_{-\infty}^{\infty} g(\tau)h(t - \tau) d\tau$$

$$g * h \iff G(f)H(f)$$

Operador...
Interacción...
Multiplicación "element-wise"

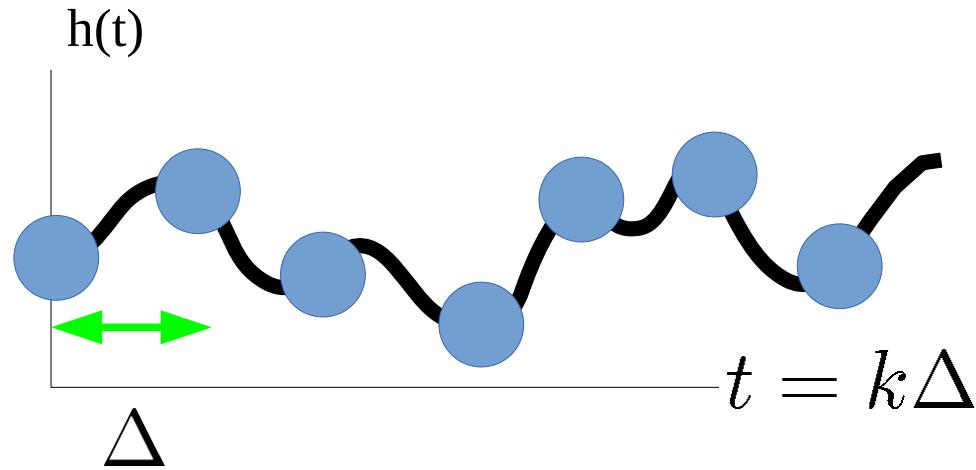
Correlation theorem

$$\text{Corr}(g, h) \equiv \int_{-\infty}^{\infty} g(\tau + t)h(\tau) d\tau$$

$$\text{Corr}(g, h) \iff G(f)H^*(f)$$

Transformada de Fourier discreta

$$h_k \equiv h(t_k), \quad t_k \equiv k\Delta, \quad k = 0, 1, 2, \dots, N - 1$$



Sampling theorem

$$H(f_n) = \int_{-\infty}^{\infty} h(t) e^{2\pi i f_n t} dt \approx \sum_{k=0}^{N-1} h_k e^{2\pi i f_n t_k} \Delta = \Delta \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

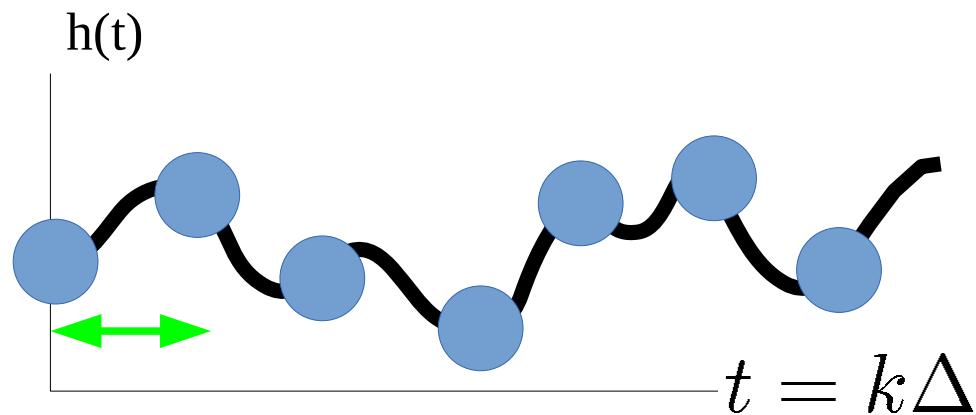
$$f_n \equiv \frac{n}{N\Delta}, \quad n = -\frac{N}{2}, \dots, \frac{N}{2}$$

Transformada
discreta

$$H_n \equiv \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

Transformada de Fourier discreta

$$h_k \equiv h(t_k), \quad t_k \equiv k\Delta, \quad k = 0, 1, 2, \dots, N - 1$$



Transformada
discreta

$$H_n \equiv \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

Anti-
Transformada

$$h_k = \frac{1}{N} \sum_{n=0}^{N-1} H_n e^{-2\pi i k n / N}$$

Se podría poner de distinta forma... con un $1/\sqrt{N}$ en ambos lados

Transformada de Fourier discreta

Transformada discreta 1D

$$H_n \equiv \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

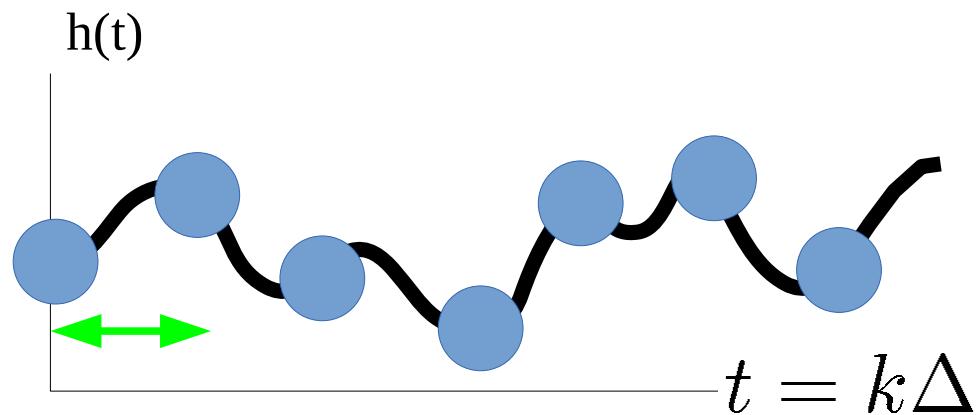
Transformada discreta 2D

$$H(n_1, n_2) \equiv \sum_{k_2=0}^{N_2-1} \sum_{k_1=0}^{N_1-1} \exp(2\pi i k_2 n_2 / N_2) \exp(2\pi i k_1 n_1 / N_1) h(k_1, k_2)$$

Transformada discreta 3D, etc

Transformada de Fourier discreta

$$h_k \equiv h(t_k), \quad t_k \equiv k\Delta, \quad k = 0, 1, 2, \dots, N - 1$$



Transformada
discreta

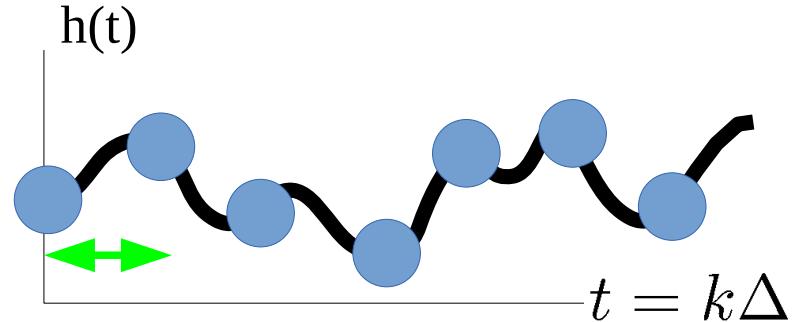
$$H_n \equiv \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

Anti-
Transformada

$$h_k = \frac{1}{N} \sum_{n=0}^{N-1} H_n e^{-2\pi i k n / N}$$

Transformada de Fourier discreta

$$h_k \equiv h(t_k), \quad t_k \equiv k\Delta, \quad k = 0, 1, 2, \dots, N - 1$$



Transformada
discreta

$$H_n \equiv \sum_{k=0}^{N-1} h_k e^{2\pi i k n / N}$$

Anti-
Transformada

$$h_k = \frac{1}{N} \sum_{n=0}^{N-1} H_n e^{-2\pi i k n / N}$$

Como calcularlas eficientemente?

Algoritmos
Secuenciales

- {
- Naïve approach [hasta ~1960]: $O(N^2)$
 - Algoritmos de *Fast Fourier transform (FFT)* [Cooley–Tukey o Gauss?]: $O(N \log_2 N) \rightarrow$ recursivo

La transformación es una transformación
lineal

Fast Fourier Transforms

Está en C y C++

La librería pública de FFT mas popular:

“La mas rápida del oeste”



Our benchmarks, performed on a variety of platforms, show that FFTW's performance is typically superior to that of other publicly available FFT software, and is even competitive with vendor-tuned codes. In contrast to vendor-tuned codes, however, FFTW's performance is portable: the same program will perform well on most architectures without modification. Hence the name, "FFTW," which stands for the somewhat whimsical title of "Fastest Fourier Transform in the West."  Anda en Intel, AMD, ...

Otra librería... pero paga...

Vendor-tuned codes:

MKL: Math Kernel Library

Intel's Math Kernel Library (MKL) is a library of optimized, math routines for science, engineering, and financial applications. Core math functions include BLAS, LAPACK, ScaLAPACK, Sparse Solvers, **Fast Fourier Transforms**, and Vector Math.

The library supports Intel and compatible processors and is available for Windows, Linux and Mac OS X operating systems.

cuFFT

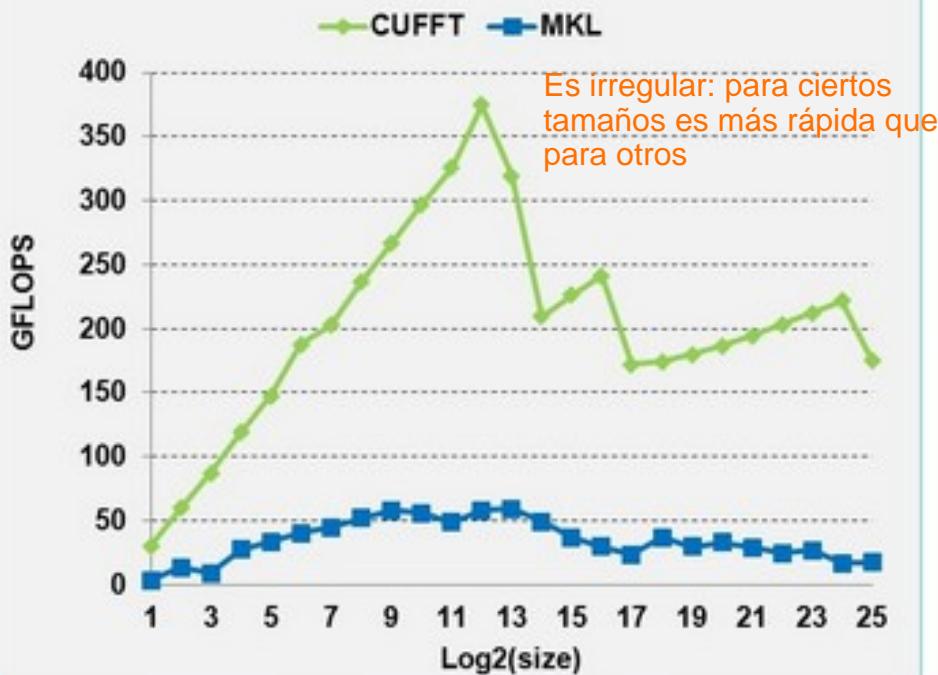
FFTS UP TO 10X FASTER THAN MKL

1D used in audio processing and as a foundation for 2D&3D FFTs

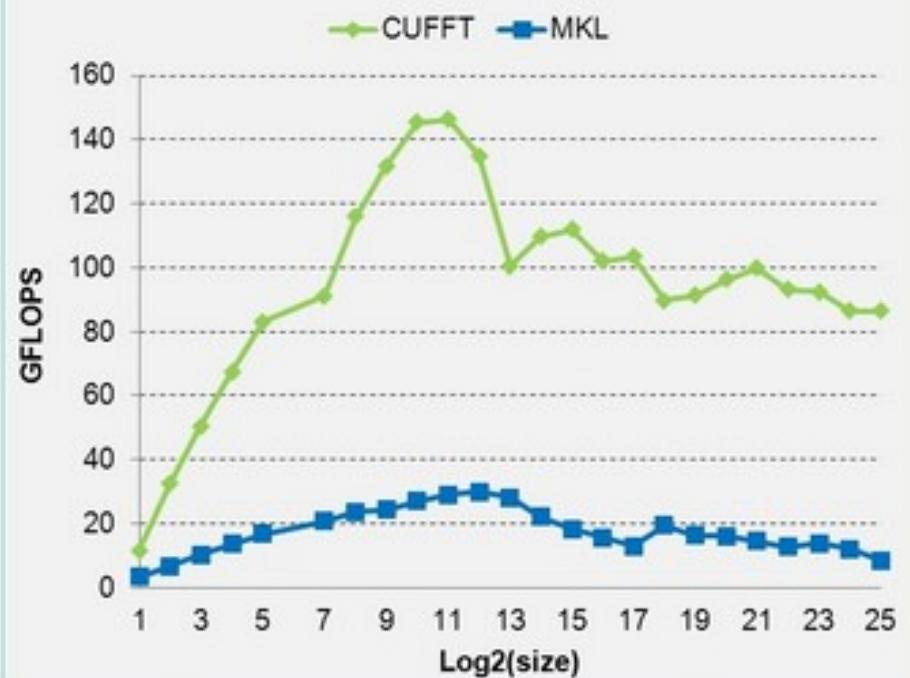
Cambia la velocidad entre simple y doble precisión, lo mejor es usar en lo posible simple precisión

Cuántas operaciones puede hacer por segundo

cuFFT Single Precision



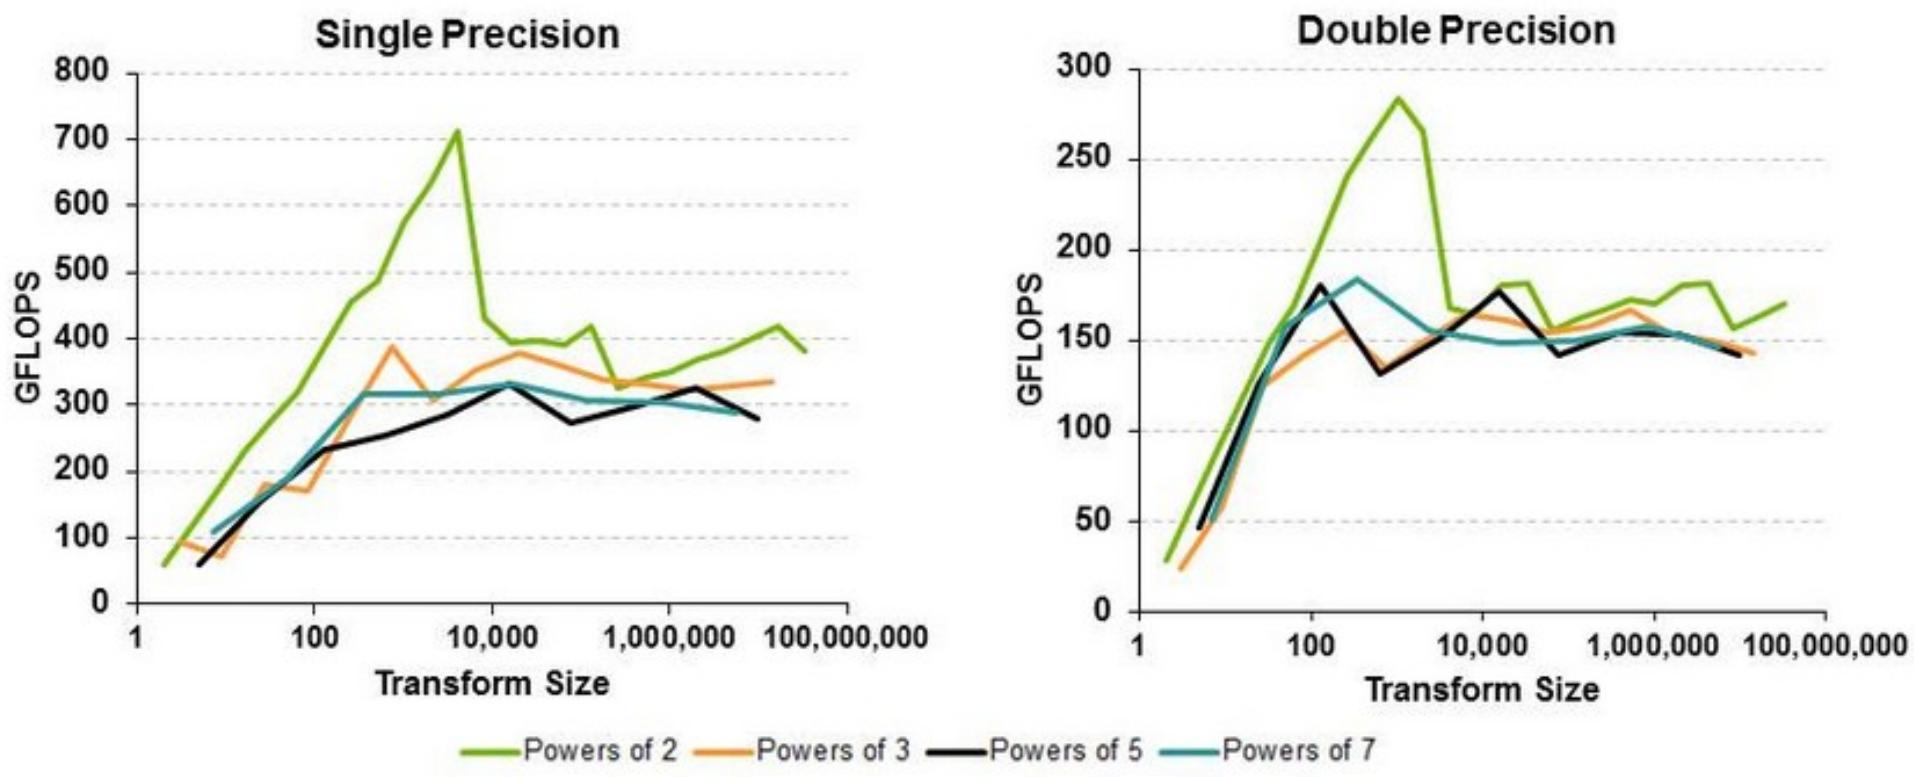
cuFFT Double Precision



- Measured on sizes that are exactly powers-of-2
- cuFFT 4.1 on Tesla M2090, ECC on
- MKL 10.2.3, TYAN FT772-B7015 Xeon x5680 Six-Core @ 3.33 GHz

- MKL 10.2.3, TYAN FT772-B7015 Xeon x5680 Six-Core @ 3.33 GHz
- Performance may vary based on OS version and motherboard configuration

CUFFT

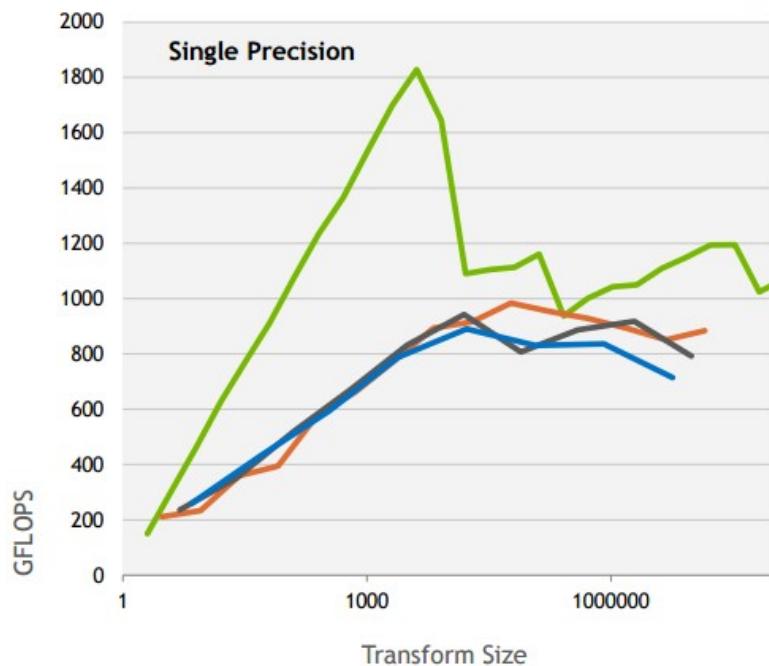


Performance may vary based on OS version and motherboard configuration

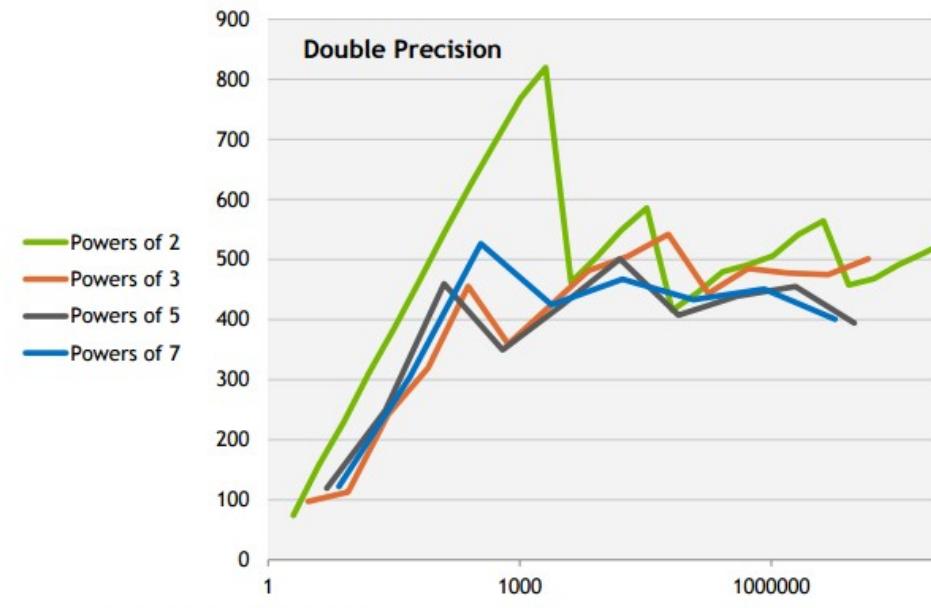
• cuFFT 6.0 on K40c, ECC ON, 28M-33M elements, input and output data on device

De 700 (K40) a 1800 (P100) GFLOPS

cuFFT: > 1800 GFLOPS SINGLE PRECISION 1D Complex, Batched FFTs



Performance may vary based on OS and software versions, and motherboard configuration



- cuFFT 8 on P100, Base clocks (r361)
- Batched transforms on 28-33M elements
- Input and output data on device
- Excludes time to create cuFFT "plans"
- Host system: Intel Xeon Haswell single-socket 16-core E5-2698 v3@ 2.3GHz, 3.6GHz Turbo
- CentOS 7.2 x86-64 with 128GB System Memory

CUFFT Samples

[simpleCUFFT - Simple CUFFT](#)

Example of using CUFFT. In this example, CUFFT is used to compute the 1D-convolution of some signal with some filter by transforming both into frequency domain, multiplying them together, and transforming the signal back to time domain.

[simpleCUFFT_2d_MGPU - SimpleCUFFT_2d_MGPU](#)

Example of using CUFFT. In this example, CUFFT is used to compute the 1D-convolution of some signal with some filter by transforming both into frequency domain, multiplying them together, and transforming the signal back to time domain on Multiple GPU.

[simpleCUFFT_callback - Simple CUFFT Callbacks](#)

Example of using CUFFT. In this example, CUFFT is used to compute the 1D-convolution of some signal with some filter by transforming both into frequency domain, multiplying them together, and transforming the signal back to time domain. The difference between this example and the Simple CUFFT example is that the multiplication step is done by the CUFFT kernel with a user-supplied CUFFT callback routine, rather than by a separate kernel call.

[simpleCUFFT_MGPU - Simple CUFFT_MGPU](#)

Example of using CUFFT. In this example, CUFFT is used to compute the 1D-convolution of some signal with some filter by transforming both into frequency domain, multiplying them together, and transforming the signal back to time domain on Multiple GPU.

Ejemplo de CUFFT

3D Complex-to-Complex Transforms

```
#define NX 64
#define NY 64
#define NZ 128
cufftHandle plan;
cufftComplex *data1, *data2;
cudaMalloc((void**)&data1, sizeof(cufftComplex)*NX*NY*NZ);
cudaMalloc((void**)&data2, sizeof(cufftComplex)*NX*NY*NZ);
/* Create a 3D FFT plan. */
cufftPlan3d(&plan, NX, NY, NZ, CUFFT_C2C);
/* Transform the first signal in place. */
cufftExecC2C(plan, data1, data1, CUFFT_FORWARD);
/* Transform the second signal using the same plan. */
cufftExecC2C(plan, data2, data2, CUFFT_FORWARD);
/* Destroy the cuFFT plan. */
cufftDestroy(plan);
cudaFree(data1); cudaFree(data2);
```

Este plan informa a la librería qué queremos hacer y la librería se encarga de buscar el mejor algoritmo

Se le informa qué tipo de transformada
~~> se va a hacer

Se está asumiendo que los datos están ordenados continuamente (que no hay stride)

No hace falta un plan por transformada, pero tiene que ser el mismo tipo de datos

Ejemplo de CUFFT

Batched transforms

```
#define NX 256
#define BATCH 10
#define RANK 1
...
{
    cufftHandle plan;
    cufftComplex *data;
    ...
    cudaMalloc((void**)&data, sizeof(cufftComplex)*NX*BATCH);

    cufftPlanMany(&plan, RANK, NX, &iembed, istride, idist,
    &oembed, ostride, odist, CUFFT_C2C, BATCH);
    ...
    cufftExecC2C(plan, data, data, CUFFT_FORWARD);
    cudaDeviceSynchronize();
    ...
    cufftDestroy(plan);
    cudaFree(data);
}
```

CUFFT y Thrust

```
#include <cufft.h>
typedef cufftDoubleReal REAL;
typedef cufftDoubleComplex COMPLEX;
// includes de thrust

int main(){
    // declaro/aloço arrays input, y su transformada, output
    thrust::device_vector<REAL> input(N);
    thrust::device_vector<COMPLEX> output(N/2+1);
    // llenar input[] con alguna señal ...

    // cast a punteros normales a memoria del device
    REAL * d_input_raw = thrust::raw_pointer_cast(&input[0]);
    COMPLEX * d_output_raw = thrust::raw_pointer_cast(&output[0]);

    cufftHandle planfft; // declara un plan ...
    cufftPlan1d(&planfft,N,CUFFT_R2C,1); // lo inicializa ...
    cufftExecR2C(planfft, d_input_raw, d_output_raw); // y lo ejecuta ...

    // copio la transformada del device al host, y la imprimo
    thrust::host_vector<COMPLEX> output_host = output;
    for(i=0; i < N/2+1; i++) cout << output_host[i].x << " " << output_host[i].y << endl;
}
```

// declaro/aloço arrays input, y su transformada, output

thrust::device_vector<REAL> input(N);

thrust::device_vector<COMPLEX> output(N/2+1);

// llenar input[] con alguna señal ...

// cast a punteros normales a memoria del device

REAL * d_input_raw = thrust::raw_pointer_cast(&input[0]);

COMPLEX * d_output_raw = thrust::raw_pointer_cast(&output[0]);

cufftHandle planfft; // *declara un plan ...*

cufftPlan1d(&planfft,N,CUFFT_R2C,1); // *lo inicializa ...*

cufftExecR2C(planfft, d_input_raw, d_output_raw); // *y lo ejecuta ...*

// copio la transformada del device al host, y la imprimo

thrust::host_vector<COMPLEX> output_host = output;

for(i=0; i < N/2+1; i++) cout << output_host[i].x << " " << output_host[i].y << endl;

Debería tener como output N elementos PERO como son complejos, realmente tengo el doble de datos. El "+1" está por las dudas si N es impar. Además, todas las frecuencias negativas están consideradas dadas las propiedades de simetría de la transformada de fourier

Ejemplo 1: Transformada 1d

- /share/apps/icnpg/clases/clase11/Simple_cufft_fftw_thrust
- **simple_cufft.cu**: transformada en CUDA
- Compilar y Correr simple_cufft:
 - **¿Es correcta la transformada?**

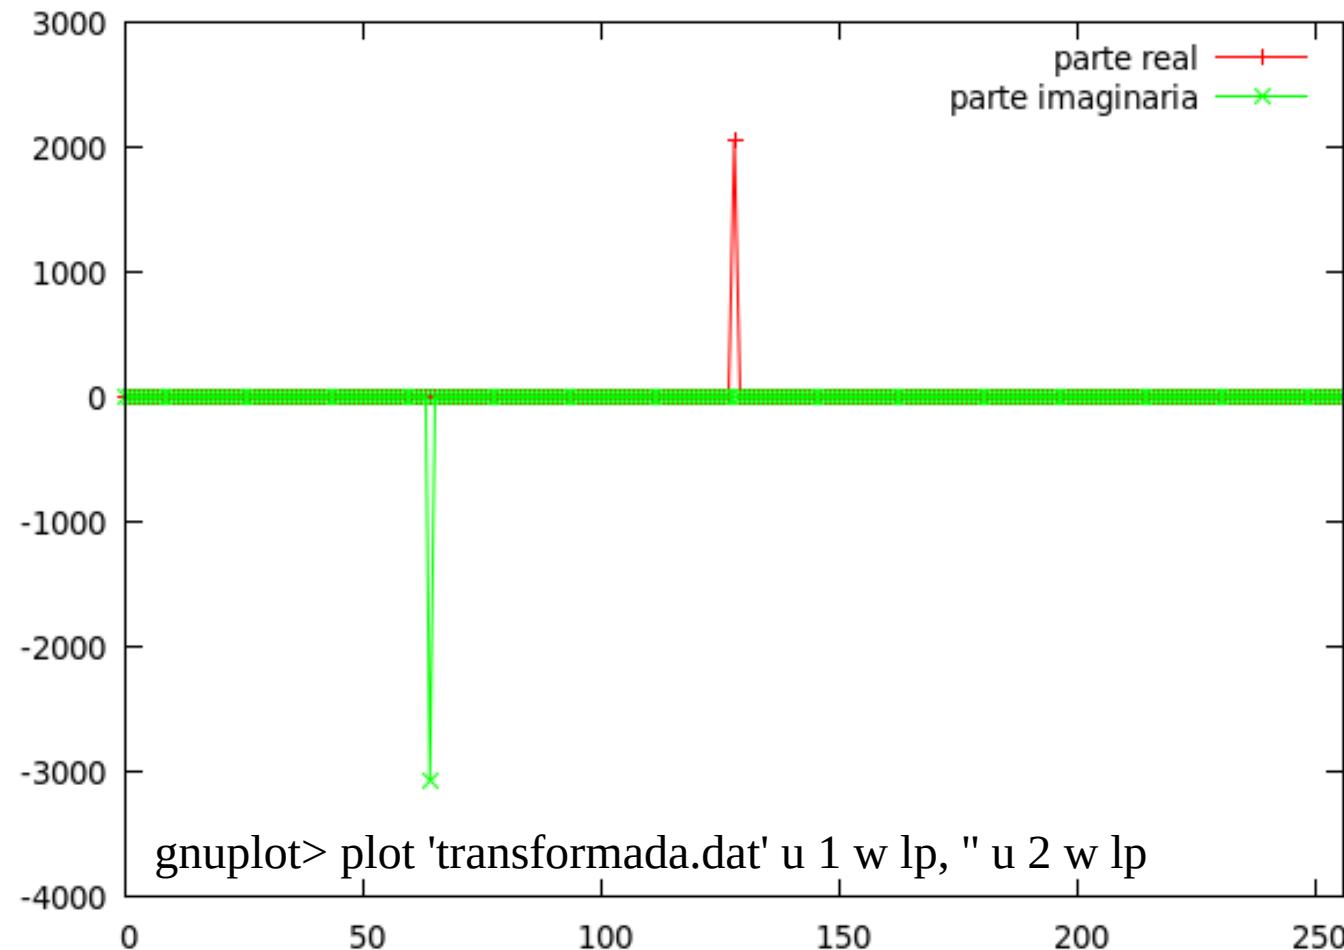
Ejemplo 1: Transformada 1d

// la transformada de este array

Input[i] = 2 A1 cos(2*M_PI*i*T1/N) + 2 A2 sin(2*M_PI*i*T2/N);

// ¿es esto? :

```
#define N 512
#define A1 4
#define A2 6
#define T1 N/4
#define T2 N/8
```



Ejemplo 1: Transformada 1d

simple_cufft.cu

```
#include <cufft.h>
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

```
#ifdef DOUBLE_PRECISION
    typedef cufftDoubleReal REAL;
    typedef cufftDoubleComplex COMPLEX;
#else
    typedef cufftReal REAL;
    typedef cufftComplex COMPLEX;
#endif
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

Estructuras de datos que vamos a necesitar

```
// Un container de thrust para guardar el input real en GPU
thrust::device_vector<REAL> D_input(N);

// toma el raw_pointer del array de input, para pasarselo a CUFFT luego
REAL *d_input = thrust::raw_pointer_cast(&D_input[0]);

// Un container de thrust para guardar el output complejo en GPU = transformada del input
int Ncomp=N/2+1;
thrust::device_vector<COMPLEX> D_output(Ncomp);

// toma el raw_pointer del array de output, para pasarselo a CUFFT luego
COMPLEX *d_output = thrust::raw_pointer_cast(&D_output[0]);
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

```
/* Parametros de la senial */
#define A1 4
#define A2 6
#define T1 N/4
#define T2 N/8
struct FillSignal
{
    __device__ __host__
    REAL operator()(unsigned tid)
    {
        // empiece con esta funcion, luego ponga cualquiera que quiera transformar...
        return A1*2.0*cosf(2*M_PI*tid*T1/(float)N) + A2*2.0*sinf(2*M_PI*tid*T2/(float)N);
    }
};
```

Llenado de la señal en el device

```
// thrust::device_vector<REAL> D_input(N);

thrust::transform(thrust::make_counting_iterator(0), thrust::make_counting_iterator(N),
D_input.begin(),FillSignal());
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

Creación de planes

```
// crea el plan de transformada de cuFFT

#ifndef DOUBLE_PRECISION

cufftHandle plan_d2z;
CUFFT_SAFE_CALL(cufftPlan1d(&plan_d2z,N,CUFFT_D2Z,1));

#else

cufftHandle plan_r2c;
CUFFT_SAFE_CALL(cufftPlan1d(&plan_r2c,N,CUFFT_R2C,1));

#endif
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

Creación de planes

```
// crea el plan de transformada de cuFFT

#ifndef DOUBLE_PRECISION

cufftHandle plan_d2z;
CUFFT_SAFE_CALL(cufftPlan1d(&plan_d2z,N,CUFFT_D2Z,1));

#else

cufftHandle plan_r2c;
CUFFT_SAFE_CALL(cufftPlan1d(&plan_r2c,N,CUFFT_R2C,1));

#endif
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

Ejecución de la transformada

```
/* ---- Start ---- */
// un timer para GPU
gpu_timer tfft;
tfft.tic();

//Transforma Fourier ejecutando el plan
#ifdef DOUBLE_PRECISION
CUFFT_SAFE_CALL(cufftExecD2Z(plan_d2z, d_input, d_output));
#else
CUFFT_SAFE_CALL(cufftExecR2C(plan_r2c, d_input, d_output));
#endif
tfft.tac();
/* ---- Stop ---- */
```

Ejemplo 1: Transformada 1d

simple_cufft.cu

Copia al host de la transformada, parte real e imaginaria

```
// declara un vector para copiar/guardar la transformada en el host:  
thrust::host_vector<COMPLEX> H_output=D_output;
```

```
#ifdef IMPRIMIR  
    ofstream transformada_out("transformada.dat");  
    for(int j = 0 ; j < Ncomp ; j++){  
        transformada_out << COMPLEX(H_output[j]).x << " " << COMPLEX(H_output[j]).y << endl;  
    }  
#endif
```

Ejemplo 1: Transformada 1d

- /share/apps/icnpg/clases/clase11/Simple_cufft_fftw_thrust
- **simple_cufft.cu**: transformada en CUDA
- Compilar y Correr simple_cufft:
 - ¿Es correcta la transformada?
- Completar template de **simple_cufft.cu**:
 - **Implemente la transformada inversa de la transformada obtenida.**
- **simple_fftw_thrust_solucion.cpp**: transformadas serial y openmp
- Comparar performances.

Guía del usuario

Mas adelante

Nuevo: útil para optimizar
- kernel fusion -

Chapter 1. Introduction.....
Chapter 2. Using the cuFFT API.....
2.1. Accessing cuFFT.....
2.2. Fourier Transform Setup.....
2.3. Fourier Transform Types.....
2.4. Data Layout.....
2.4.1. FFTW Compatibility Mode.....
2.5. Multidimensional Transforms.....
2.6. Advanced Data Layout.....
► 2.7. Streamed cuFFT Transforms.....
► 2.8. Multiple GPU cuFFT Transforms.....
2.8.1. Plan Specification and Work Areas.....
2.8.2. Helper Functions.....
2.8.3. Multiple GPU 2D and 3D Transforms on Permuted Input.....
2.8.4. Supported Functionality.....
2.9. cuFFT Callback Routines.....
2.9.1. Overview of the cuFFT Callback Routine Feature.....
2.9.2. Specifying Load and Store Callback Routines.....
2.9.3. Callback Routine Function Details.....
2.9.4. Coding Considerations for the cuFFT Callback Routine Feature
2.10. Thread Safety.....
2.11. Static Library and Callback Support.....
2.12. Accuracy and Performance.....

Callbacks

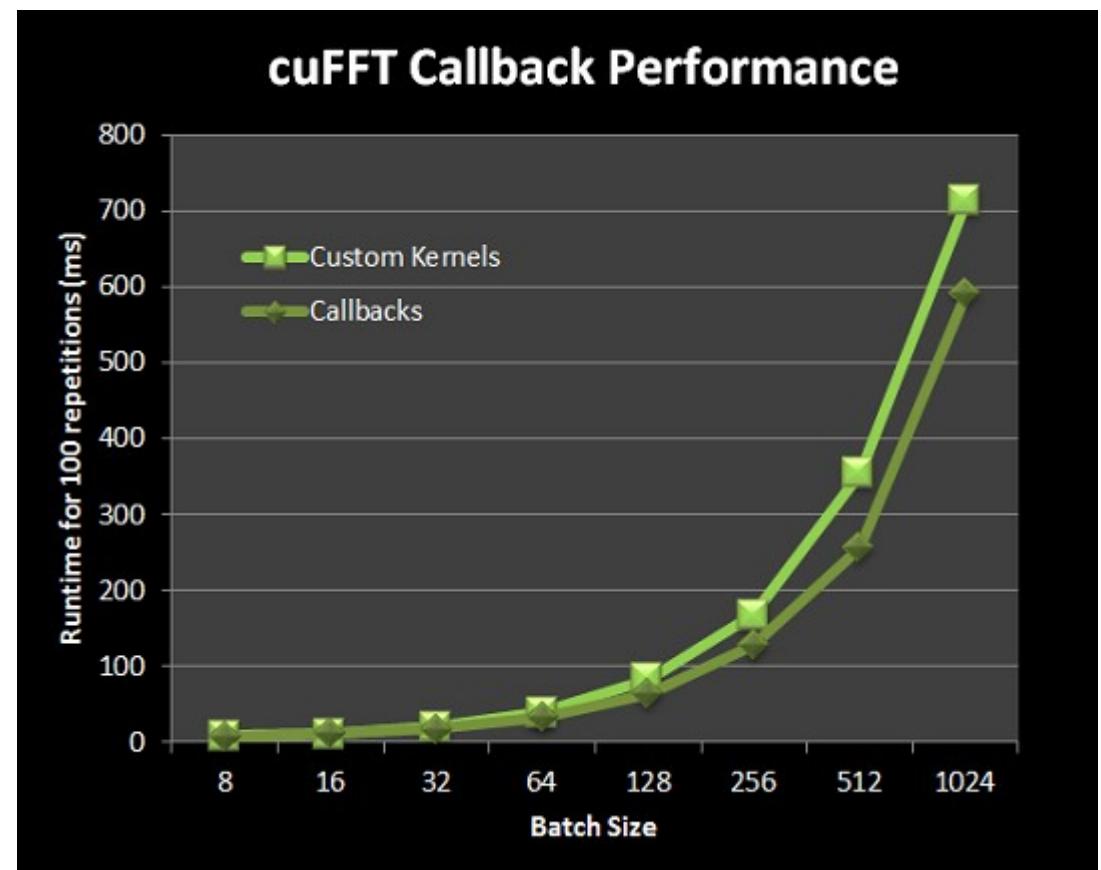
Before CUDA 6.5: 3 kernels, 3 memory roundtrips



With CUDA 6.5: 1 kernel, 1 memory roundtrip



- <https://devblogs.nvidia.com/cuda-pro-tip-use-cufft-callbacks-custom-data-processing/>



Guía del usuario

Chapter 7. FFTW INTERFACE TO CUFFT

NVIDIA provides FFTW3 interfaces to the cuFFT library. This allows applications using FFTW to use NVIDIA GPUs with minimal modifications to program source code. To use the interface first do the following two steps

- ▶ It is recommended that you replace the include file **fftw3.h** with **cufftw.h**
- ▶ Instead of linking with the double/single precision libraries such as **fftw3/fftw3f** libraries, link with both the cuFFT and cuFFTW libraries
- ▶ Ensure the search path includes the directory containing **cuda_runtime_api.h**

After an application is working using the FFTW3 interface, users may want to modify their code to move data to and from the GPU and use the routines documented in the [FFTW Conversion Guide](#) for the best performance.

Referencias

- CUFFT Library User Guide NVIDIA.
- CUDA Libraries SDK Code Samples (simpleCUFFT, PDEs, etc).
- FFTW home page: <http://www.fftw.org/> [interface parecida para CPU + openmp, mpi].
- Numerical Recipes in C [algoritmo FFT explicado].
- Libros de Numerical PDEs.... [como discretizar las PDEs?]

Ejemplo 2: Transformada 1d en cupy

CuPy

Overview Installation User Guide API Reference Contribution Guide Upgrade Guide License

Search the docs ...

The N-dimensional array (`ndarray`)
Universal functions (`cupy.ufunc`)
Routines (NumPy)
Array creation routines
Array manipulation routines
Binary operations
Data type routines
Discrete Fourier Transform (`cupy.fft`)
`cupy.fft.fft`
`cupy.fft.ifft`
`cupy.fft.ifft2`
`cupy.fft.ifft2`
`cupy.fft.ifftn`
`cupy.fft.ifftn`
`cupy.fft.rfft`
`cupy.fft.irfft`
`cupy.fft.rfft2`
`cupy.fft.irfft2`
`cupy.fft.rfftn`
`cupy.fft.irfftn`
`cupy.fft.hfft`
`cupy.fft.ihfft`
`cupy.fft.rfftfreq`
`cupy.fft.rfftfreq`
`cupy.fft.fftshift`

Discrete Fourier Transform (`cupy.fft`)

Hint
NumPy API Reference: Discrete Fourier Transform (numpy.fft)

See also
Discrete Fourier transforms (`cupyx.scipy.fft`)

Standard FFTs

`fft(a[, n, axis, norm])` Compute the one-dimensional FFT.

`ifft(a[, n, axis, norm])` Compute the one-dimensional inverse FFT.

`fft2(a[, s, axes, norm])` Compute the two-dimensional FFT.

`ifft2(a[, s, axes, norm])` Compute the two-dimensional inverse FFT.

`fftn(a[, s, axes, norm])` Compute the N-dimensional FFT.

`ifftn(a[, s, axes, norm])` Compute the N-dimensional inverse FFT.

Real FFTs

`rfft(a[, n, axis, norm])` Compute the one-dimensional FFT for real input.

`irfft(a[, n, axis, norm])` Compute the one-dimensional inverse FFT for real input.

<https://docs.cupy.dev/en/stable/reference/fft.html>

Ejemplo 3: convolución 2d en CUDA C++

- /share/apps/icnpg/clases/Clases_cufft/Convolucion2d_cufft
- Convoluciona o filtra un arreglo bidimensional usando cufft:
 - Transformo filtro (analiticamente)
 - Transformo imagen (cufft)
 - Multiplico las transformadas complejas
 - Antitransformo el producto y normalizo debidamente...
- ¿Funciona?
- ¿Para qué sirve?

Ejemplo 4: operadores diferenciales 2d

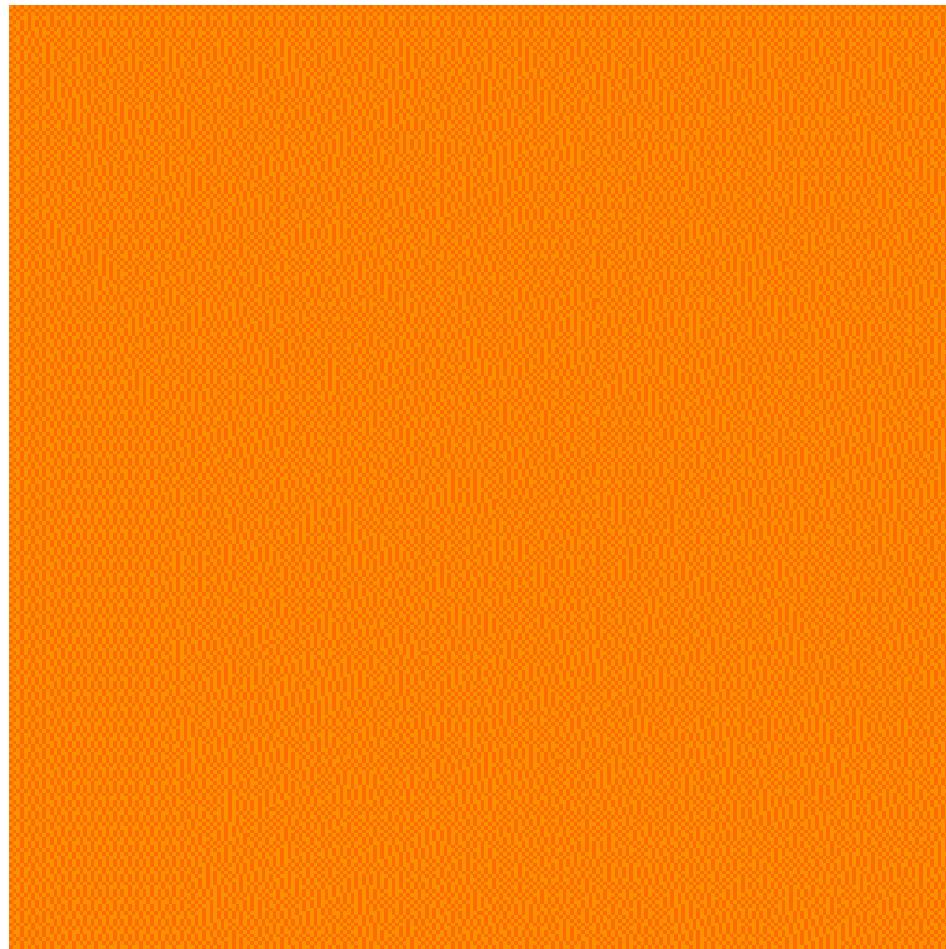
Ecuación “del calor” en dos dimensiones con fuentes arbitrarias en el espacio de Fourier

$$\frac{\partial}{\partial t} u(x, y, t) = c \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) u(x, y, t) + \rho(x, y, t)$$

Ecuación de difusión con fuentes arbitrarias

Una configuración de fuentes divertida

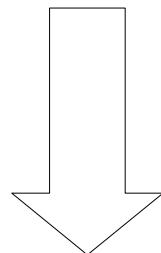
“Placa
Metálica”
[Condiciones
de Contorno
Periódicas]



En las diapositivas esto era
una animación

Ecuación “del calor” en dos dimensiones con fuentes arbitrarias en el espacio de Fourier

$$\frac{\partial}{\partial t} u(x, y, t) = c \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) u(x, y, t) + \rho(x, y, t)$$



Transformando Fourier

Modos desacoplados

$$\frac{\partial}{\partial t} \tilde{u}(q_x, q_y, t) = -c (q_x^2 + q_y^2) \tilde{u}(q_x, q_y, t) + \tilde{\rho}(q_x, q_y, t)$$

$$\tilde{u}(q_x, q_y, t = 0) = 0 \rightarrow$$

$$\tilde{u}(q_x, q_y, t) = \int_0^t dt' e^{-c(q_x^2 + q_y^2)(t-t')} \tilde{\rho}(q_x, q_y, t')$$

Solución analítica!

Ecuación de difusión con fuentes arbitrarias

$$\frac{\partial}{\partial t} u(x, y, t) = c \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) u(x, y, t) + \rho(x, y, t)$$

Diferencias finitas espaciales

$$\begin{cases} \delta_y^2 u_{i,l} = u_{j,l+1} - 2u_{j,l} + u_{j,l-1} \\ \delta_x^2 u_{i,l} = u_{j+1,l} - 2u_{j,l} + u_{j-1,l} \end{cases}$$

Esquema explícito (sin fuente)

$$u_{j,l}(t + \Delta t) = u_{j,l}(t) + \frac{c\Delta t}{\Delta^2} (\delta_x^2 u_{j,l}(t) + \delta_y^2 u_{j,l}(t)), \quad \frac{c\Delta t}{\Delta^2} \leq \frac{1}{2}$$

Esquema implícito [Crank-Nicholson]

$$u_{j,l}(t + \Delta t) = u_{j,l}(t) +$$

$$\frac{c\Delta t}{2\Delta^2} (\delta_x^2 u_{j,l}(t) + \delta_x^2 u_{j,l}(t + \Delta t) + \delta_y^2 u_{j,l}(t) + \delta_y^2 u_{j,l}(t + \Delta t))$$

Ecuación de difusión con fuentes arbitrarias

$$\frac{\partial}{\partial t} u(x, y, t) = c \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) u(x, y, t) + \rho(x, y, t)$$

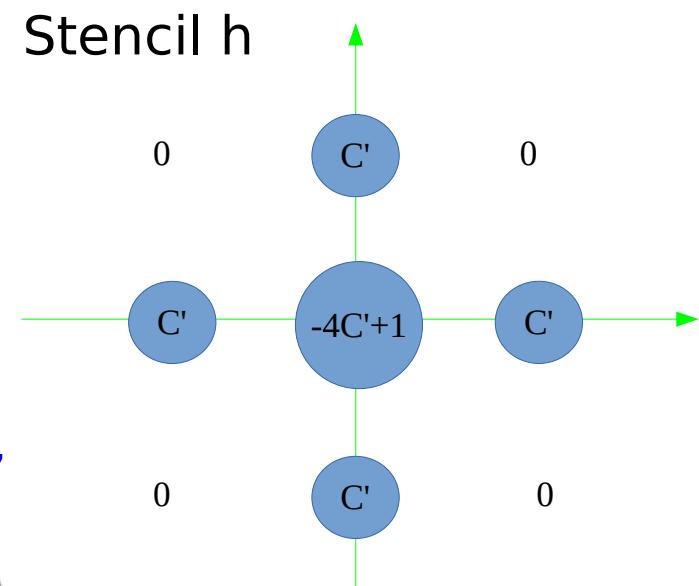
Diferencias finitas

$$u_{i,j}(t + \Delta t) \approx u_{i,j}(t) +$$

$$\frac{c\Delta t}{\Delta^2} (u_{i+1,j}(t) + u_{i-1,j}(t) + u_{i,j+1}(t) + u_{i,j-1}(t) - 4u_{i,j}(t)) + \Delta t \rho_{i,j}(t)$$

$$u_{i,j}(t + \Delta t) \approx \sum_{a,b} u_{i+a,j+b} h_{a,b} + \delta t \rho_{i,j}(t)$$

Convolucion en 2D

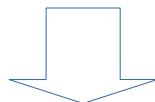


Se puede resolver en el espacio real (ej: con CUSP),
Pero lo haremos en espacio de Fourier (cuFFT)

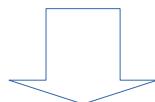
$$g * h \iff G(f)H(f)$$

Ecuación de difusión con fuentes arbitrarias

$$u_{i,j}(t + \Delta t) \approx \sum_{a,b} u_{i+a,j+b} h_{a,b} + \delta t \rho_{i,j}(t)$$

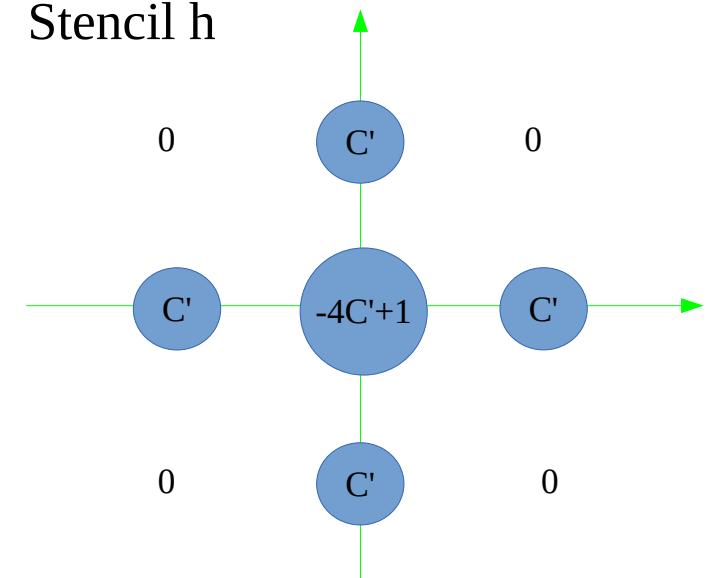


$$g * h \iff G(f)H(f)$$



Todo desacoplado :-)

Stencil h



$$\tilde{u}_{n,m}(t + \Delta t) \approx \tilde{u}_{n,m}(t) \tilde{h}_{n,m} + \delta t \tilde{\rho}_{n,m}(t)$$

$$\tilde{h}_{n,m} = \sum_{a,b} e^{2\pi i(na/L_x + mb/L_y)} h_{a,b}$$

$$\tilde{h}_{n,m} = (-4C' + 1) + 2C'[\cos(2\pi(n/L_x)) + \cos(2\pi(m/L_y))]$$

Aplicar el stencil o hacer la convolución 2D en el espacio real se reduce a multiplicar por un campo escalar constante la transformada

Ecuación de difusión con fuentes arbitrarias

Esto se calcula una sola vez al principio

$$\tilde{h}_{n,m} = (-4C' + 1) + 2C'[\cos(2\pi(n/L_x)) + \cos(2\pi(m/L_y))]$$

Loop temporal

```
for(t=0;t<TRUN;t++)  
{
```

$$\tilde{u}_{n,m}(t + \Delta t) \approx \tilde{u}_{n,m}(t)\tilde{h}_{n,m} + \delta t \tilde{\rho}_{n,m}(t), \quad \forall n, m$$

```
}
```

Euler Step
Trivialmente
Paralelizable

Ecuación de difusión con fuentes arbitrarias

$$\tilde{u}_{n,m}(t + \Delta t) \approx \tilde{u}_{n,m}(t)\tilde{h}_{n,m} + \delta t\tilde{\rho}_{n,m}(t), \quad \forall n, m$$

Euler Step
Trivialmente
Paralelizable

```
// Paso de Euler en el espacio de Fourier
__global__ void euler_step
(cufftComplex *a, const cufftComplex *b,
const float *qqmodule, int nx, int ny)
{
    // idx and idy, the location of the element in the original NxN array
    int idx = blockIdx.x*blockDim.x+threadIdx.x;
    int idy = blockIdx.y*blockDim.y+threadIdx.y;
    if ( idx < nx && idy < ny){
        int index = idx*ny + idy;
        const float fac=- (qqmodule[index]);
        a[index].x = (1.f+C_BETA*fac*DT)*a[index].x + DT*b[index].x;
        a[index].y = (1.f+C_BETA*fac*DT)*a[index].y + DT*b[index].y;
    }
}
```

$\tilde{u}_{n,m}(t + \Delta t)$

$\tilde{h}_{n,m}$

$\tilde{u}_{n,m}(t)$

$\tilde{\rho}_{n,m}(t)$

Conviene trabajar
Con grillas 2D!

Ecuación de difusión con fuentes arbitrarias

```
for(t=0;t<TRUN;t++)  
{
```

$$\tilde{u}_{n,m}(t) \rightarrow u_{i,j}(t)$$

...

$$\rho_{i,j}(t) \rightarrow \tilde{\rho}_{n,m}(t)$$

...

$$\tilde{u}_{n,m}(t + \Delta t) \approx \tilde{u}_{n,m}(t) \tilde{h}_{n,m} + \delta t \tilde{\rho}_{n,m}(t), \quad \forall n, m$$

```
}
```

Para dibujar (cada tanto) el campo en espacio real → **CUFFT (inv)** → memcpy

Las fuentes son dadas instantáneamente en el espacio real, pero las necesitamos en Fourier → **CUFFT**

Kernel Paso de Euler

Ecuación de difusión con fuentes arbitrarias

```
for(t=0;t<TRUN;t++)  
{
```

$$\tilde{u}_{n,m}(t) \rightarrow u_{i,j}(t)$$

...

$$\rho_{i,j}(t) \rightarrow \tilde{\rho}_{n,m}(t)$$

...

$$\tilde{u}_{n,m}(t + \Delta t) \approx \tilde{u}_{n,m}(t)\tilde{h}_{n,m} + \delta t\tilde{\rho}_{n,m}(t), \quad \forall n, m$$

```
}
```

Para dibujar (cada tanto) el campo en espacio real → **CUFFT** → memcpy

Las fuentes son dadas instantáneamente en el espacio real, pero las necesitamos en Fourier → **CUFFT**

Kernel Paso de Euler

Hands-on: Ecuación de difusión con fuentes arbitrarias

```
gpu_timer cronometro;
cronometro.tic();
for(int i=0;i<T_RUN+1; i++){                                Loop temporal principal

    /* Get data each T_DATA steps*/
    if(i%T_DATA==0) {
        /* Take the scalar field back from Fourier space and normalize */
        T.AntitransformFromFourierSpace();
        T.Normalize();

        T.CpyDeviceToHost(); // To be used in a moment for printing.

        /* Print frame for visualization */
        sprintf(filename, "frame%d.ppm", 100000000+i);
        ifstream file(filename);
        T.PrintPicture(filename,i);
    }

    float phase=2.f*M_PI*i*4.f/T_RUN;
    float x0=0.2*sin(phase);float y0=0.2*cos(phase);           float intensidad=0.1;
    T.InitParticular(1, LX,LY,4,x0,y0,intensidad); // perfil de fuentes instantaneo
    T.TransformForceToFourierSpace(); // transforma perfil de fuentes para Euler Step
}

/* Integration step in Fourier Space*/
T.EulerStep();
```

main.cu

En este problema:
Solo para hacer imágenes

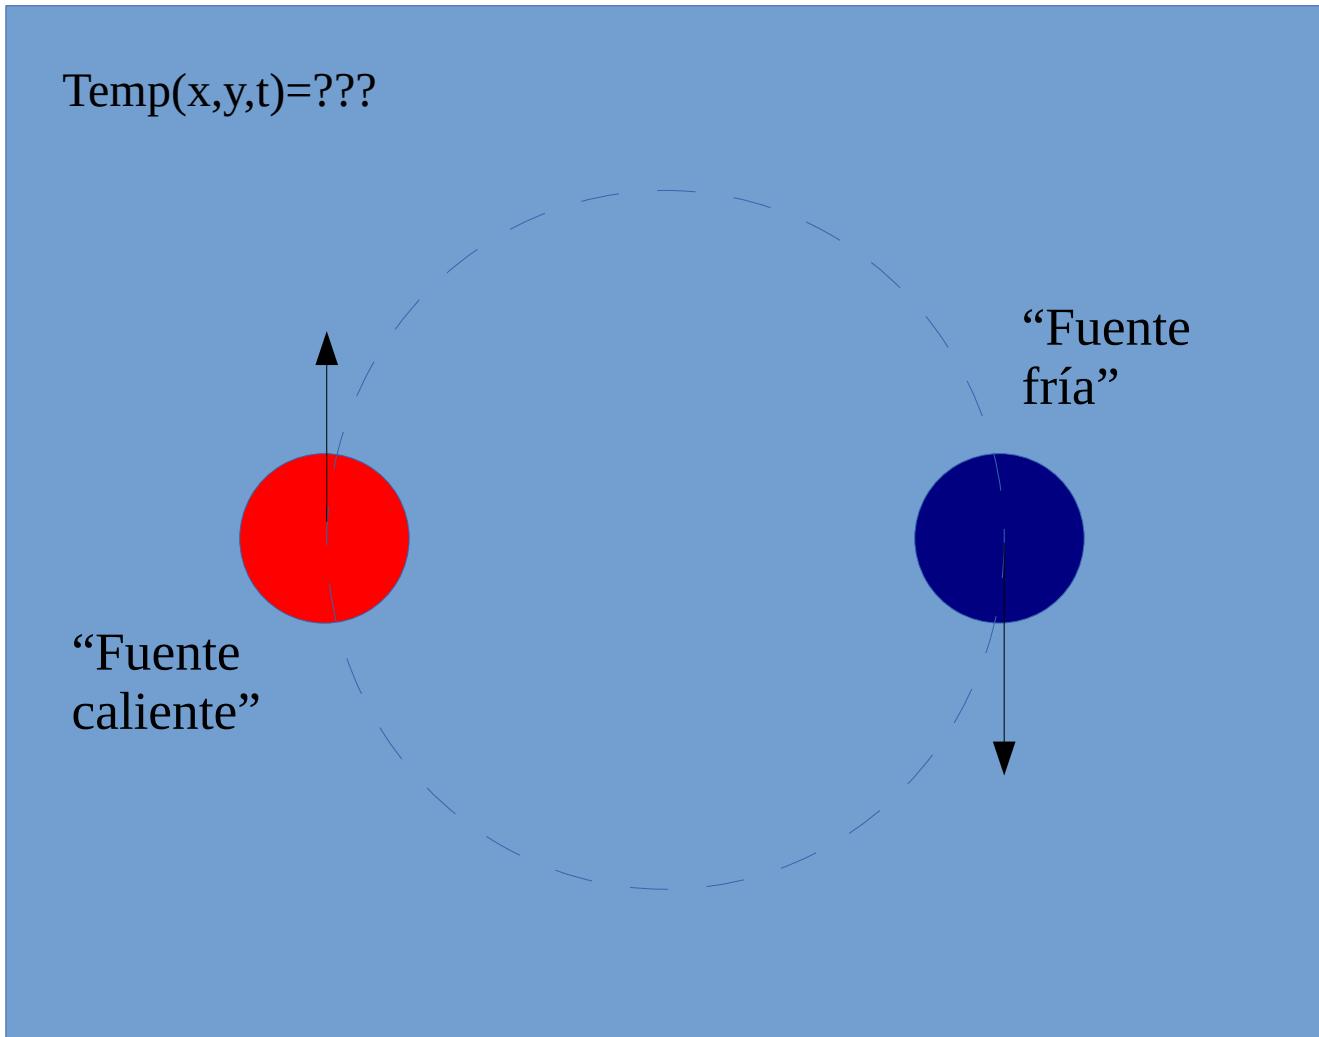
Controla fuentes
en espacio real...

Un paso de tiempo en el
Espacio de Fourier

Hands-on: Ecuación de difusión con fuentes arbitrarias

Una configuración de fuentes divertida

“Placa
Metálica”
[Condiciones
de Contorno
Periódicas]



Hands-on!

Ecuación de difusión con fuentes arbitrarias

- /share/apps/icnpg/clases/Clases_cufft/Heat/
- make
- qsub jobGPU
- *Escribe snapshots.ppm → película.*

Referencias

- CUFFT Library User Guide NVIDIA.
- CUDA Libraries SDK Code Samples (simpleCUFFT, PDEs, etc).
- FFTW home page: <http://www.fftw.org/> [interface parecida para CPU + openmp, mpi].
- Numerical Recipes in C [algoritmo FFT explicado].
- Libros de Numerical PDEs.... [como discretizar las PDEs?]

Un ejemplo concreto de aplicación científica

- **Formación de patrones de magnetización en films ferromagnéticos:**
Matemáticamente: Ecuación diferencial parcial no local en función del tiempo para un campo escalar bidimensional.

Un modelo simple para la magnetización en films ferromagnéticos 2D

$$\phi(\mathbf{r}, t) = \phi(x, y, t)$$

Magnetización coarse-grained (“soft spin”),
o parámetro de orden bidimensional

Landau approach: modelar la energía libre en función del parámetro de orden

$$H_l = \alpha_0 \int d\mathbf{r} \left(-\frac{\phi(\mathbf{r})^2}{2} + \frac{\phi(\mathbf{r})^4}{4} \right) - h_0 \int d\mathbf{r} \phi(\mathbf{r})$$

Dos mínimos locales, y un campo
Que propicia uno de los dos

$$H_{rig} = \beta_0 \int d\mathbf{r} \frac{|\nabla \phi(\mathbf{r})|^2}{2}$$

Un precio energético por “doblar”
La magnetización

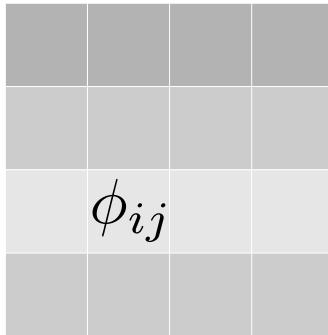
$$H_{dip} = \gamma_0 \int d\mathbf{r} d\mathbf{r}' \phi(\mathbf{r}) \phi(\mathbf{r}') G(\mathbf{r}, \mathbf{r}')$$

Interacciones dipolares de largo alcance
 $G(\mathbf{r}, \mathbf{r}') \sim 1/|\mathbf{r} - \mathbf{r}'|^3$

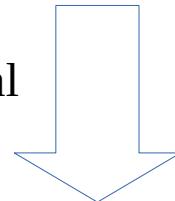
$$\frac{\partial \phi(\mathbf{r})}{\partial t} = -\lambda \frac{\delta(H_l + H_{rig} + H_{dip})}{\delta \phi(\mathbf{r})} = -\lambda \left(\alpha_0(-\phi + \phi^3) - h_0 - \beta_0 \nabla^2 \phi + \gamma_0 \int d\mathbf{r}' G(|\mathbf{r} - \mathbf{r}'|) \phi(\mathbf{r}') \right)$$

Método Naive

$$\frac{\partial \phi(\mathbf{r})}{\partial t} = -\lambda \left(\alpha_0(-\phi + \phi^3) - h_0 - \beta_0 \nabla^2 \phi + \gamma_0 \int d\mathbf{r}' G(|\mathbf{r} - \mathbf{r}'|) \phi(\mathbf{r}') \right)$$

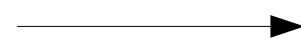


Discretización en espacio real



$$\phi_{ij}^{\text{new}} = \phi_{ij} - \lambda \delta t \left(\alpha(-\phi_{ij} + \phi_{ij}^3) - h_0 - \beta(\phi_{i+1,j} + \phi_{i-1,j} + \phi_{i,j+1} + \phi_{i,j-1} - 4\phi_{ij}) + \gamma \sum_{kl} G_{ij,kl} \phi_{kl} \right)$$

Naive: $O(N^*N)$



$\sim N$ operaciones

$\sim N^2$ operaciones

G es densa...

G es densa ... pero tiene tanta información?

G es densa ... pero tiene un montón de simetrías!

$$G_{ij,kl} = G_{kl,ij}, \quad G_{(i+p)(j+q),(k+p)(l+q)} = G_{ij,kl}$$

Transformación de Fourier

$$\phi(\mathbf{r}) = \phi(x, y)$$

Magnetización coarse-grained (“soft spin”),
o parámetro de orden bidimensional

$$\frac{\partial \phi(\mathbf{r})}{\partial t} = -\lambda \left(\alpha_0(-\phi + \phi^3) - h_0 - \beta_0 \nabla^2 \phi + \gamma_0 \int d\mathbf{r}' G(|\mathbf{r} - \mathbf{r}'|) \phi(\mathbf{r}') \right)$$

- **Dificultad:** ya no solo a primeros vecinos, sino todos los puntos del espacio 2D están acoplados

Pero, en espacio de Fourier:

$$\frac{\partial \phi_{\mathbf{k}}}{\partial t} = -\lambda [\alpha_0(-\phi + \boxed{\phi^3})|_{\mathbf{k}} - h_0 \delta(\mathbf{k}) + (\beta_0 k^2 + \gamma_0 G_{\mathbf{k}}) \phi_{\mathbf{k}}]$$

Si no fuera por este término, en el espacio de Fourier
quedarían $L_x * L_y$ ecuaciones desacopladas!!.

- Mas simple ir al espacio de Fourier para la interacción de largo alcance y volver al espacio real para calcular término cubico.
- También es eficiente si $N=Lx Ly$ es grande: N^2 vs $2 N \log N$



Solución numérica: método pseudo-spectral

$$\phi = \phi(x, y)$$

Magnetización coarse-grained (“soft spin”),
o parámetro de orden bidimensional

$$\frac{\partial \phi_{\mathbf{k}}}{\partial t} = -\lambda[\alpha_0(-\phi + \phi^3)|_{\mathbf{k}} - h_0\delta(\mathbf{k}) + (\beta_0 k^2 + \gamma_0 G_{\mathbf{k}})\phi_{\mathbf{k}}]$$

Semi-implicit Euler step:

$$\frac{\phi_{\mathbf{k}}^{t+\delta t} - \phi_{\mathbf{k}}^t}{\delta t} = [\alpha(\phi - \phi^3)|_{\mathbf{k}}^t + h_0\delta(\mathbf{k}) - \gamma G_{\mathbf{k}}\phi_{\mathbf{k}}^t - \beta k^2 \phi_{\mathbf{k}}^{t+\delta t}]$$

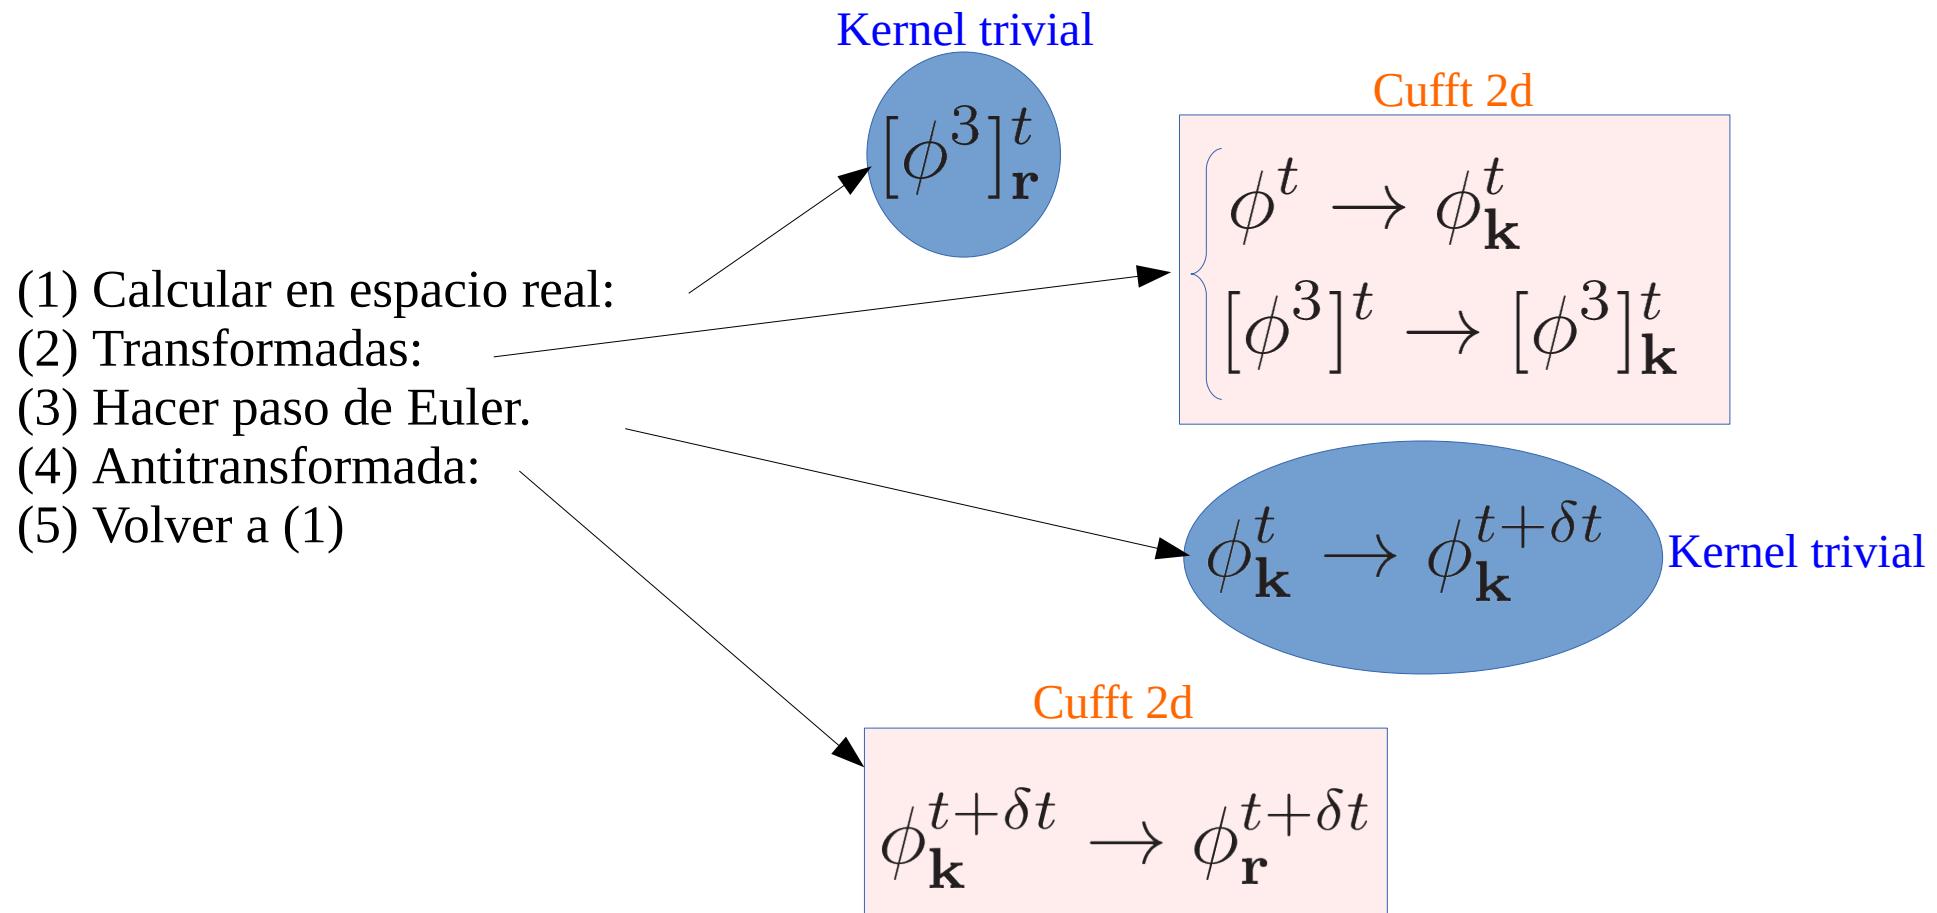
estabilidad
↓

$$\rightarrow \phi_{\mathbf{k}}^{t+\delta t} = \frac{\phi_{\mathbf{k}}^t + \delta t[\alpha(\phi - \phi^3)|_{\mathbf{k}}^t + h_0\delta(\mathbf{k}) - \gamma G_{\mathbf{k}}\phi_{\mathbf{k}}^t]}{1 + \beta k^2 \delta t}$$

Algoritmo

Update:

$$\phi_{\mathbf{k}}^{t+\delta t} = \frac{\phi_{\mathbf{k}}^t + \delta t [\alpha(\phi - \phi^3)|_{\mathbf{k}}^t + h_0 \delta(\mathbf{k}) - \gamma G_{\mathbf{k}} \phi_{\mathbf{k}}^t]}{1 + \beta k^2 \delta t}$$



Implementación

$$\frac{\partial \phi(\mathbf{r})}{\partial t} = -\lambda \left(\alpha(-\phi + \phi^3) - h_0 - \beta \nabla^2 \phi + \gamma \int d\mathbf{r}' G(|\mathbf{r} - \mathbf{r}'|) \phi(\mathbf{r}') \right)$$

Lineal pero No-local

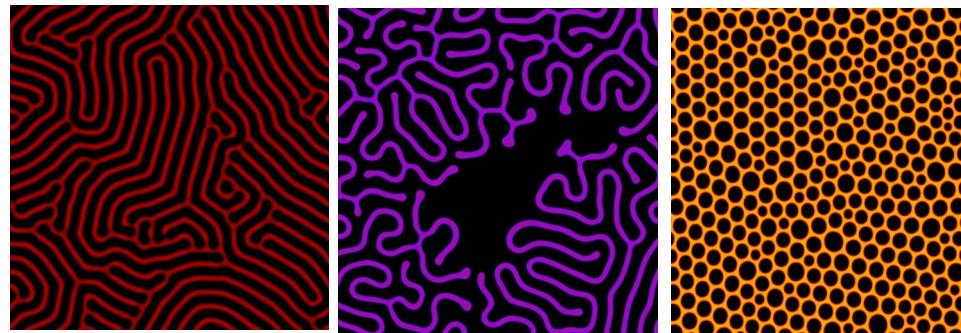
↓

$$\phi_{\mathbf{k}}^{t+\delta t} = \frac{\phi_{\mathbf{k}}^t + \delta t [\alpha(\phi - \phi^3)|_{\mathbf{k}}^t + h_0 \delta(\mathbf{k}) - \gamma G_{\mathbf{k}} \phi_{\mathbf{k}}^t]}{1 + \beta k^2 \delta t}$$

Discretización + Transformada de Fourier

Local y lineal

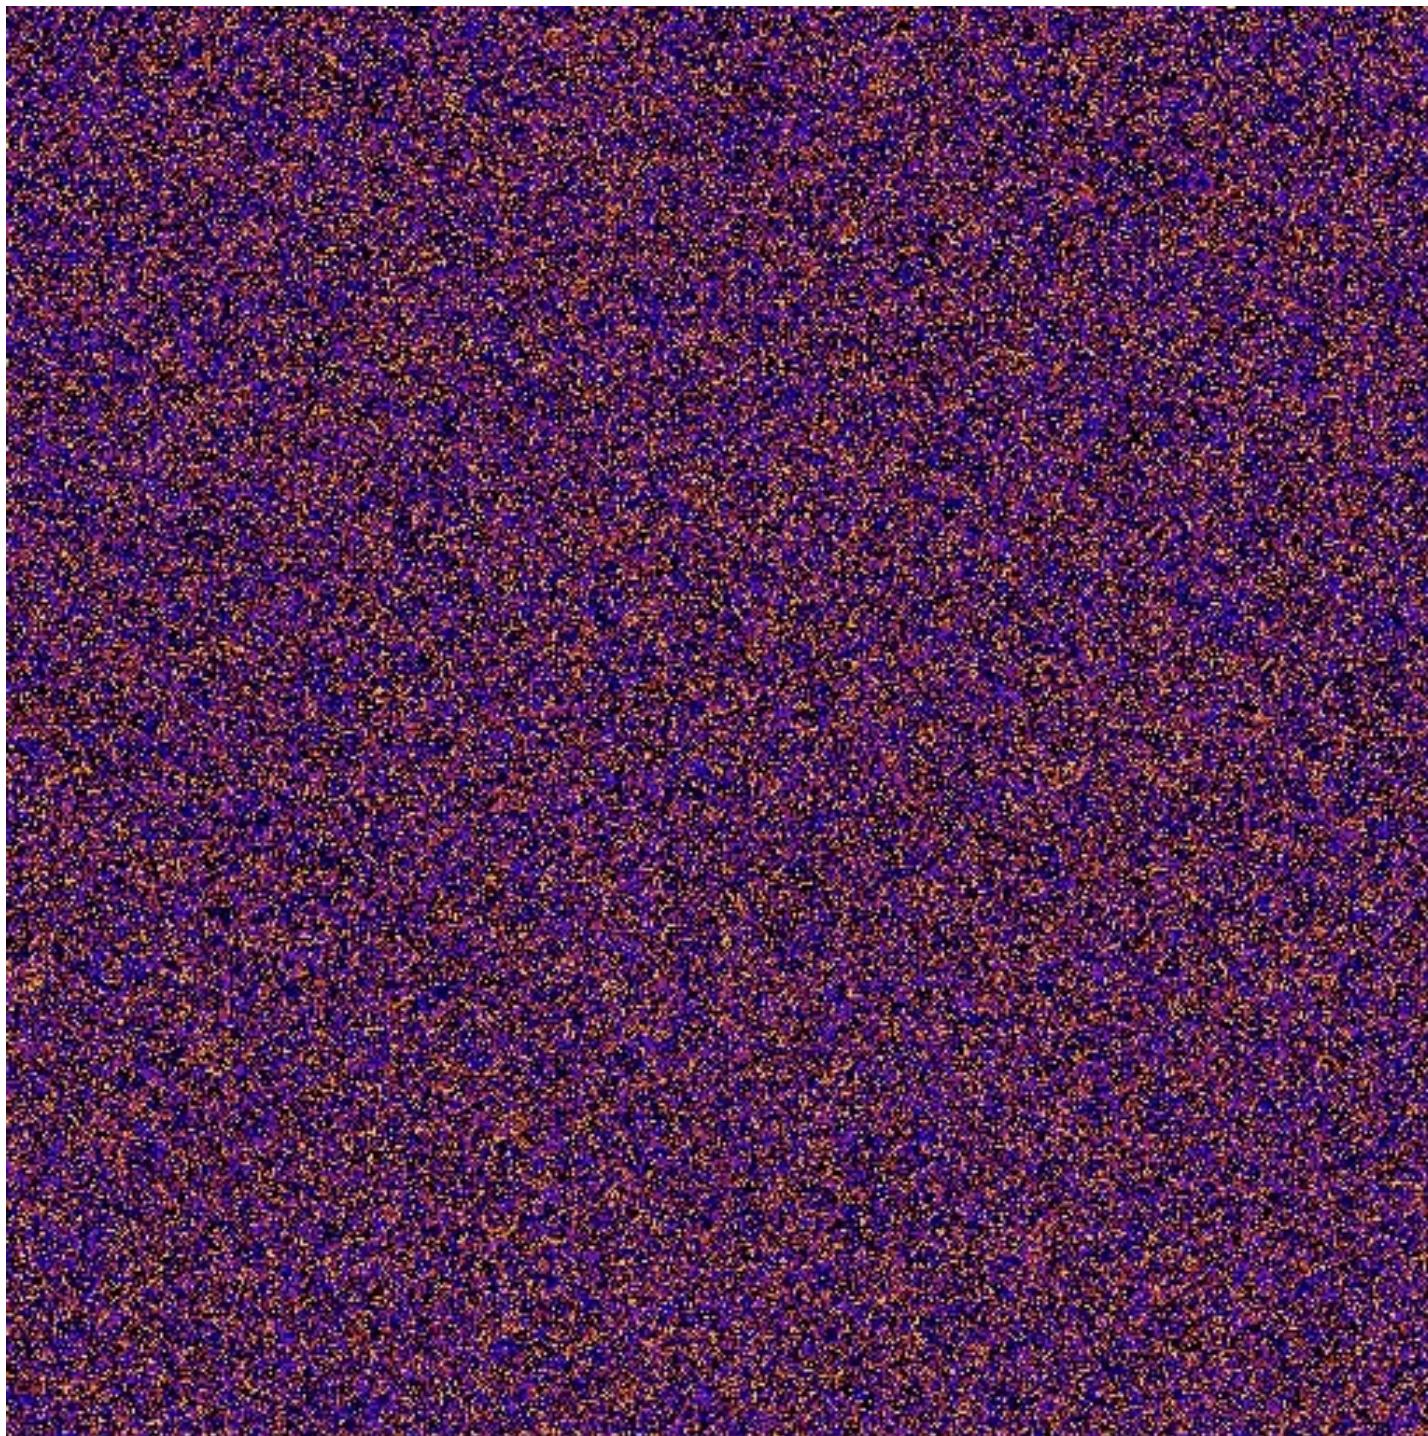
Requiere transformar/antitransformar matrices en cada paso de tiempo....

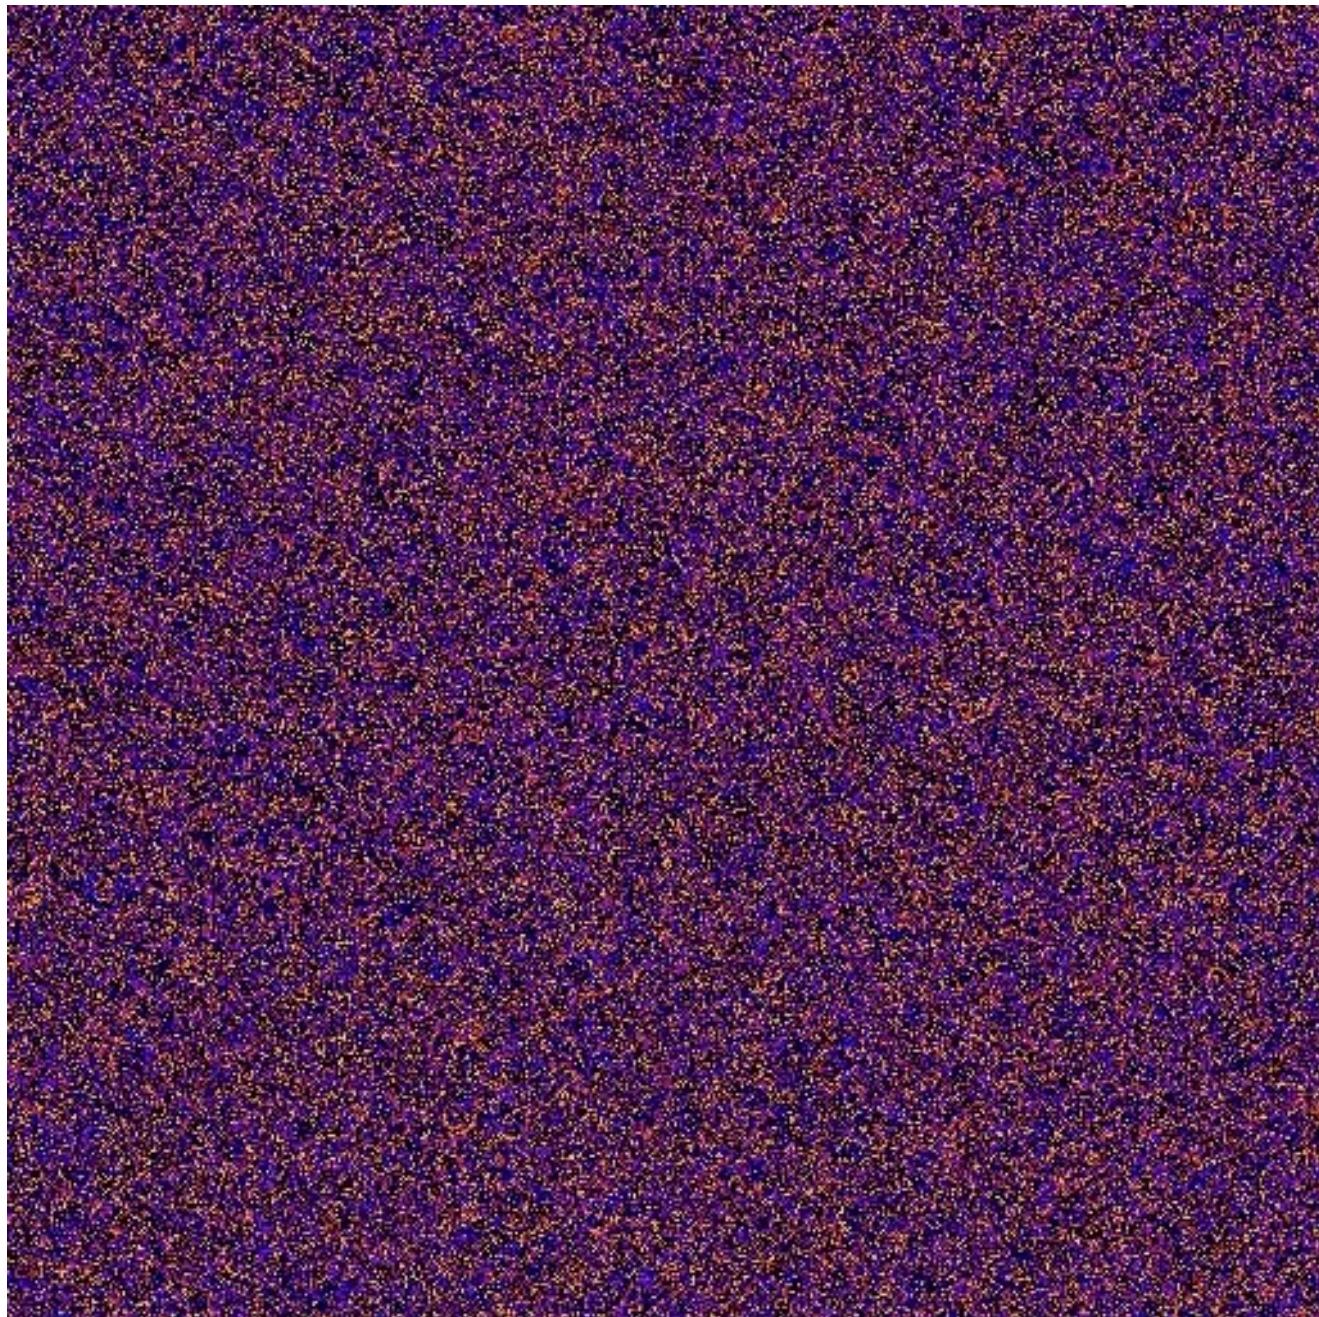


Naive:
 $O(N^2)$

PseudoEspectral:
 $O(N \log(N))$



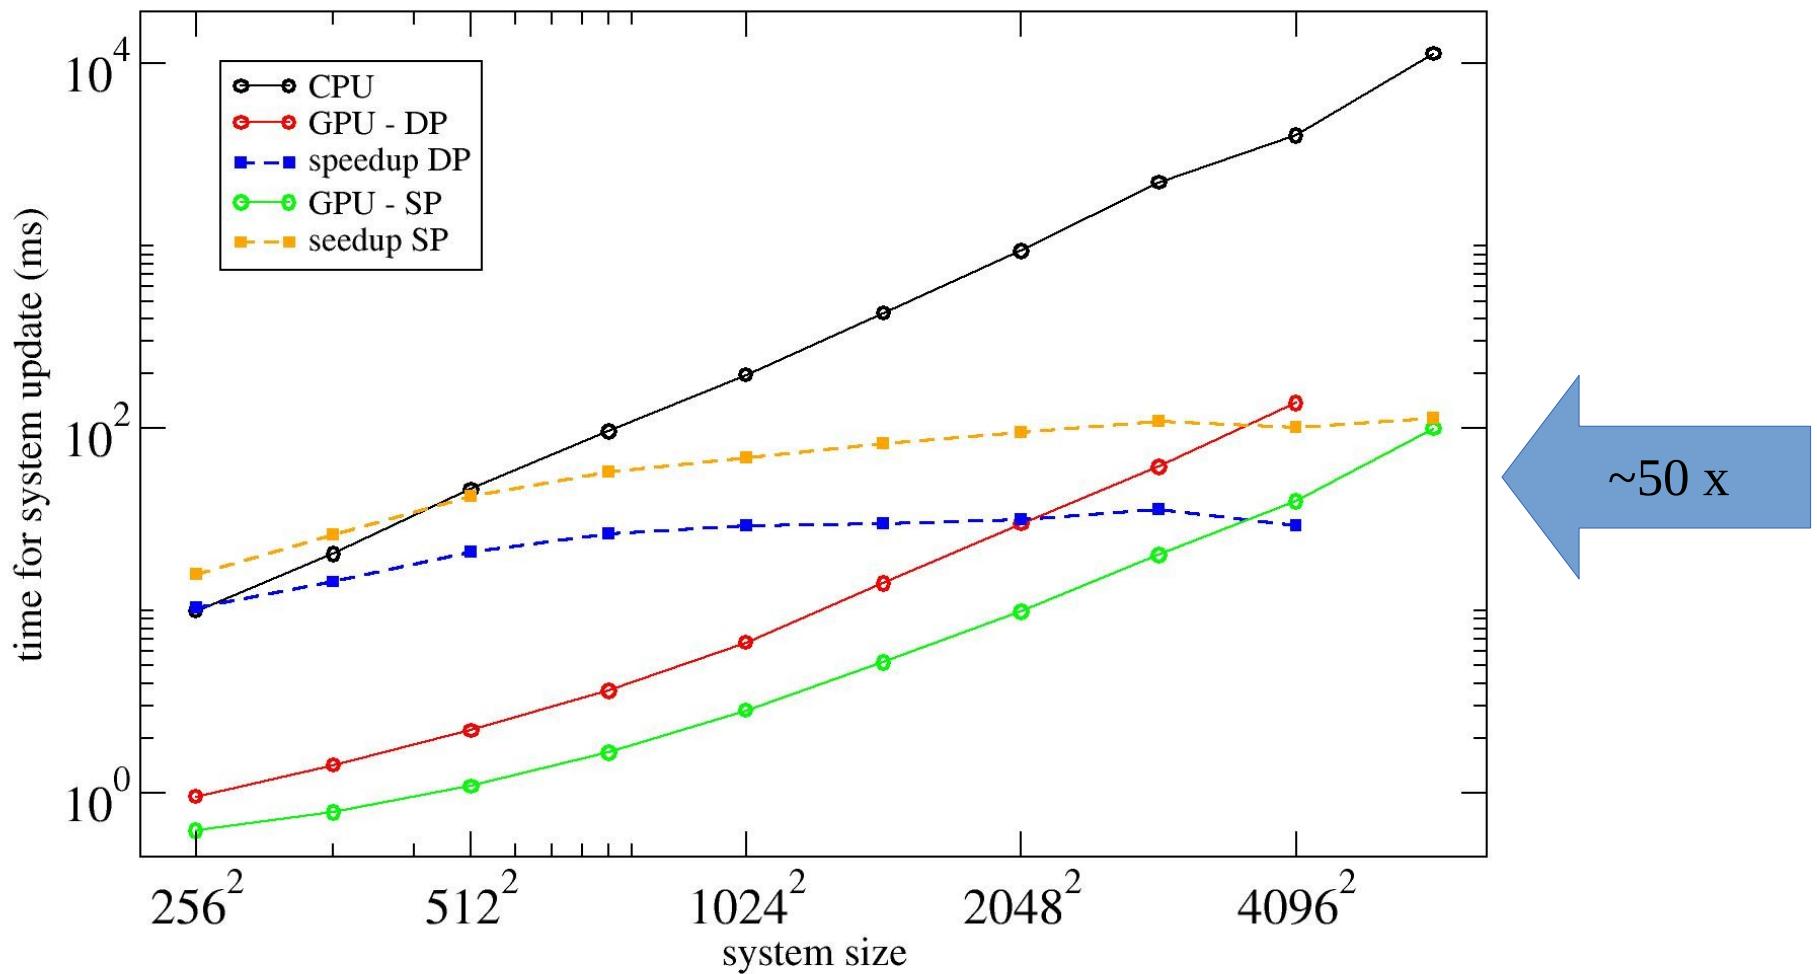






Performance

$$\phi_{\mathbf{k}}^{t+\delta t} = \frac{\phi_{\mathbf{k}}^t + \delta t [\alpha(\phi - \phi^3)|_{\mathbf{k}}^t + h_0 \delta(\mathbf{k}) - \gamma G_{\mathbf{k}} \phi_{\mathbf{k}}^t]}{1 + \beta k^2 \delta t}$$



GTX 470 vs 1 thread AMD Phenom II